

# CSE231: Operating Systems

## Assignment 3

**Name:** Mihir Chaturvedi  
**Roll no.:** 2019061  
**Section:** B  
**Branch:** CSE

---

In this assignment, we are required to modify the inbuilt CFS scheduler of the Linux kernel and allow for real-time guarantees to take precedence over the default selection method (using the RB-tree and vruntime). Following are the list of modifications that have been implemented.

### Addition of `rtnice` Attribute

This is a new attribute of the `sched_entity` attribute of a task/process that will provide information of the real-time guarantee required by the process. We add this in the new attribute `sched_entity` struct, and initialize the value of `rtnice` for a forked process to 0 in the `__sched_fork()` function.

### Changes to the Scheduler

Here, the main changes are made to the selection of the process to next execute. In the `entity_before()` function, we prioritize the `sched_entity` that has a lower `rtnice` value, and return a binary 1 or 0 depending on the comparison. This check is done before comparing the `vruntimes`, and is only done if either `sched_entities` have a positive `rtnice` value.

Next, we need to update the `rtnice` value of the `sched_entity` once it has run on the CPU. We do this in the `update_curr()` function. We reduce the `rtnice` value of the 'curr'ent entity by ``delta_exec``, the time that the process ran on the CPU, and equate it to zero if ``delta_exec`` is greater than the `rtnice` value. Here, we also increase the `vruntime` of the of the entity, since once it's back on the RB-tree (if at all) with a zero-`rtnice` value, then the

comparisons will be made using `vruntime`, which should be accurately reflected to ensure fairness.

## Addition of System Call

To modify the `rtnice` value of the `sched_entity` of a process from the userspace, we must also create a system call that does so. We create a new system call, name it `'rtnice_mod'`, and assign it the numerical ID 440.

The system call accepts two inputs/arguments from the user: the PID for the process, and the new `rtnice` value (in seconds) for the process. Since the `rtnice` attribute of the `sched_entity` is an unsigned long long that holds the values as nanoseconds, we convert the user-inputted seconds into nanoseconds.

## Testing

To test this, I have created a file that first forks the parent process into 'n' child processes. Each child process has a default `rtnice` value of 0 is made to do a long computation, after which it outputs the time taken.

Then, we fork the parent process 'n' times again, but give each new child process a positive, increasing `rtnice` value. The time taken to complete the computation is printed.

## Error Handling

- **EINVAL:** This error will be returned when an invalid PID number has been passed, or a negative `rtnice` value.
- **ESRCH:** This error will be returned when the requested PID does not correspond to any currently running process, and so the `'task_struct'` for the requested PID is NULL.

## Sample

The following is a sample out of the test file:

```
rtnice value step-size: 100
```

```
Number of forks/children: 5
```

```
With rtnice=0:
```

```
Child process 1; rtnice=0; duration=3.327755 seconds
```

```
Child process 3; rtnice=0; duration=3.675423 seconds
```

```
Child process 2; rtnice=0; duration=3.743091 seconds
```

```
Child process 5; rtnice=0; duration=3.744749 seconds
```

```
Child process 4; rtnice=0; duration=3.783028 seconds
```

```
With increasing rtnice:
```

```
Child process 2; rtnice=100; duration=2.979624 seconds
```

```
Child process 5; rtnice=400; duration=2.984560 seconds
```

```
Child process 3; rtnice=200; duration=2.985612 seconds
```

```
Child process 4; rtnice=300; duration=2.979846 seconds
```

```
Child process 1; rtnice=0; duration=5.641314 seconds
```